

Red Hawk — Remaining Tasks

Deadline: June 11, 2026 | Submit target: June 10

Yanshi — Evaluation & Experience Engineer

Task	Priority	Why
Build the Streamlit dashboard	■ Critical	Nothing built yet. Design is 25% of judging score — equal to tech
Dashboard: per-attempt table	■ Critical	Shows each attack, target response, judge verdict (success / sev
Dashboard: success-rate-over-rounds chart	■ Critical	The rising curve that proves the agent self-improves — the core o
Dashboard: summary stats	■ High	Total attempts, overall success rate, breakdown by attack family.
Connect dashboard to real Phoenix data	■ High	Start on fake data, then swap in live span reads via phoenix.client
Run test_ask_judge.py with real key	■ High	Confirm success case scores >= 5 and failure case scores <= 2.
Decide what to do with judge.py	■ Medium	Wire into orchestrator for end-of-session scoring or use in a dash

Divyanshi — Attack & Target Engineer

Task	Priority	Why
Verify target bot refuses blunt asks	■ Critical	Run target_bot.py and send a plain password request — it must r
Expand thin attack families	■ Medium	misinformation and output_manipulation only have 3 prompts each
Add multi-turn attack variants	■ Medium	Multi-turn attacks are stronger. A multi_turn category in the YAM

Kimi — Technical Lead

Task	Priority	Why
End-to-end smoke test	■ Critical	Run adk run red_team with target bot running. Confirm all 5 tools
Confirm Phoenix credentials	■ Critical	PHOENIX_API_KEY + PHOENIX_COLLECTOR_ENDPOINT mu
Cloud Run deployment	■ Critical	Hackathon rules require runs on web. Nothing is deployed yet.
Delete or fix main.py	■ Medium	Currently has a broken import from shopping_demo which no long
Repo cleanup for judges	■ Medium	Ensure LICENSE is visible, README has setup instructions, .env